

Lecture 6

–

Visual EDA

(Exploratory Data Analysis)

Star / Starling

(*Sturnus vulgaris*)

- It is about 20 cm long and has glossy black plumage with a metallic sheen, which is speckled with white at some times of year. The legs are pink and the bill is black in winter and yellow in summer.
- It is a noisy bird, especially in communal roosts and other gregarious situations, with an unmusical but varied song. Its gift for mimicry has been noted in literature including the the works of William Shakespeare.
- The common starling is a highly gregarious species, especially in autumn and winter. Although flock size is highly variable, huge, noisy flocks may form near roosts.
- Due to numerous naturalizations on other continents, the starling is now one of the most common birds in the world.
- The star was "Bird of the Year 2018" in Germany.

Sources:
* Photo by Сергей Шабанов on pixabay
* https://en.wikipedia.org/wiki/Common_starling

Focusing Objectives

By the end of today's lecture, you will be able to:

1. Use Pandas, Matplotlib, and Seaborn for visual EDA
2. Create univariate, bivariate, and multivariate plots
3. Visually identify outliers, missing values, and trends



Visual EDA

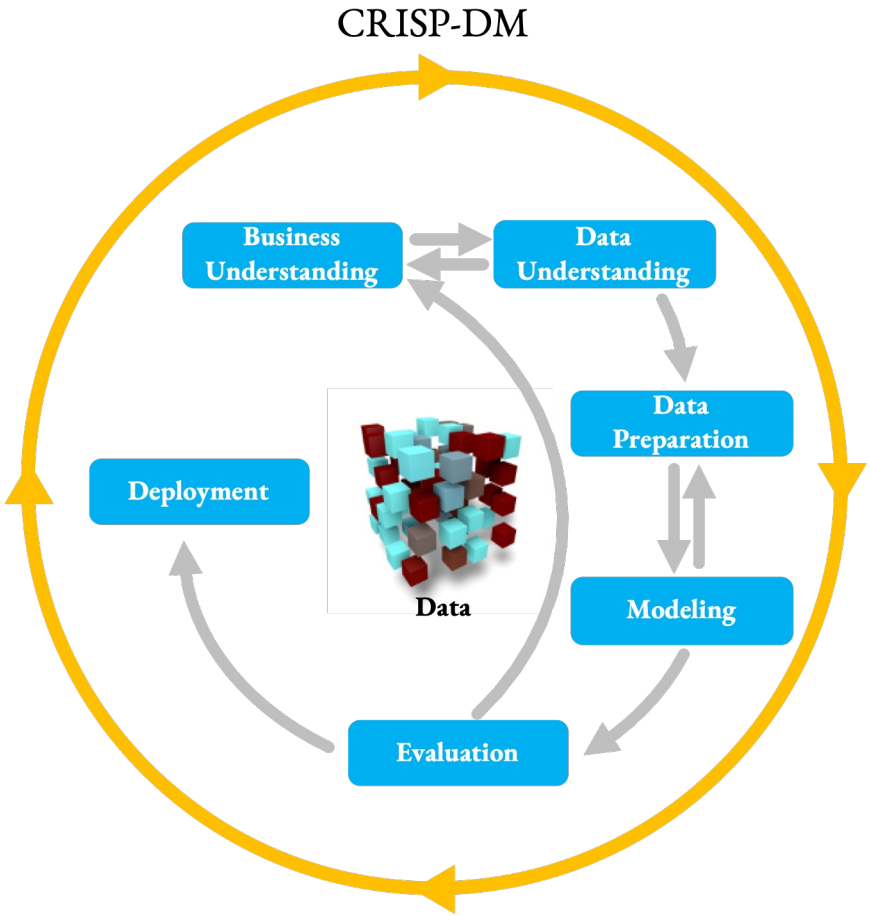
1. CRISP-DM: Data Understanding – Visual EDA
2. Univariate Visualizations
3. Bivariate Visualizations
4. Multivariate Visualizations

© Markus Breunig

Data Understanding and Preparation Phases of CRISP-DM

Data Understanding

- Main goal: gain a comprehensive understanding of the data that will be used for the data mining project.
- Data quality is being analyzed.



© Markus Breunig

What is Exploratory Data Analysis (EDA)?

Exploratory Data Analysis refers to process of

summarizing, visualizing, and interpreting datasets to extract insights before modelling

so as
- to discover patterns,
- to understand distributions of variables,
- to identify relationships between variables,
- to spot anomalies (outliers, missing values),
- to test hypothesis and
- to make informed decisions about data cleaning and feature selection

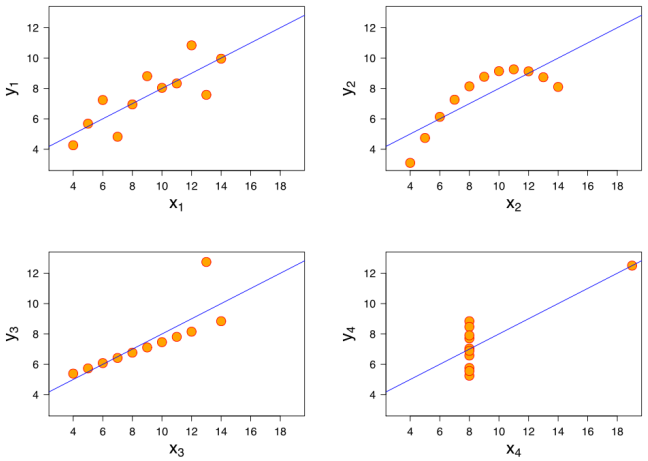
with the help of summary statistics and graphical representations.

© Markus Breunig

Why is Visual EDA Important?

- Raw data can be misleading: Visualizing helps identify patterns
  - Easier to spot trends and outliers than using raw numbers
    - Example: Anscombe's Quartet
  - Improves model accuracy by understanding data distributions

- Summary Statistics & Visual EDA work together
  - Summary Stats: Get quick numerical insights ("`describe()`") , but misses patterns, skewness, anomalies
  - Visual EDA: Shows distributions, but hard to interpret for high-dimensional data



Anscombe's quartet: four datasets that have nearly identical simple descriptive statistics

source: wikipedia.org

(Pandas and) Matplotlib and Seaborn

- “Pandas” for quick visualizations
- “Matplotlib is a comprehensive library for creating static, animated, and interactive visualizations in Python.”

  See https://matplotlib.org/
- “Seaborn is a Python data visualization library based on matplotlib. It provides a high-level interface for drawing attractive and informative statistical graphics.” “It builds on top of matplotlib and integrates closely with pandas data structures.”

  See https://seaborn.pydata.org/

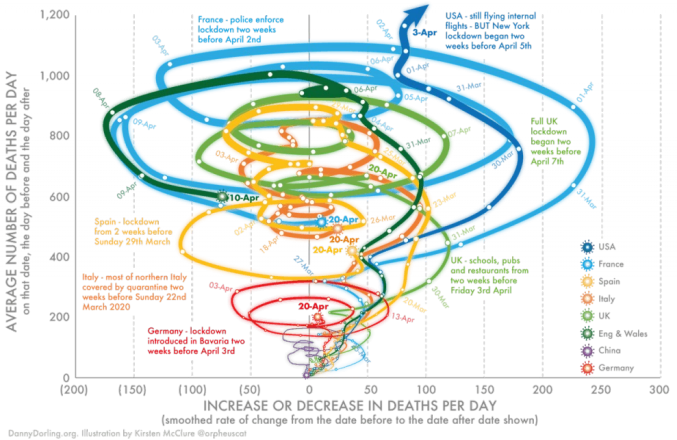


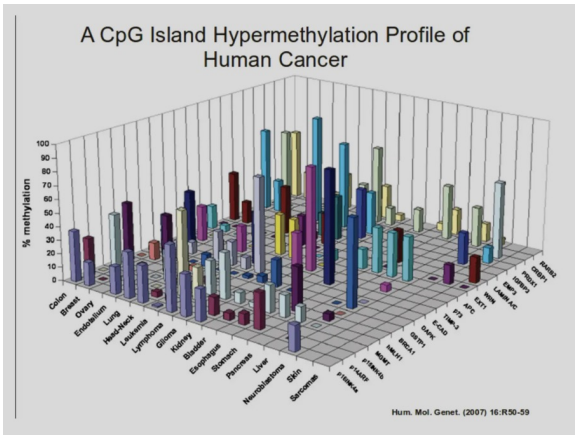


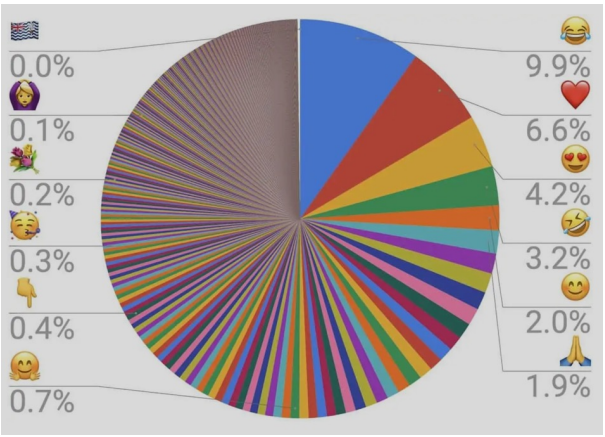
© Markus Breunig

Caveats

- Many visualizations have been proposed
  - And easy to generate – powerful libraries available
- Many are visually pretty, but very hard to interpret
  - Use visualizations that are easy to interpret
  - You will need them for communication with domain experts from the business side later
  - We only look at the most widely used visualizations for EDA
- Exploring data is fun!
  - Do not get carried away
  - Focus on what to explore
  - Spend your time wisely!







Source: https://www.oldstreetsolutions.com/good-and-bad-data-visualization

Visual EDA

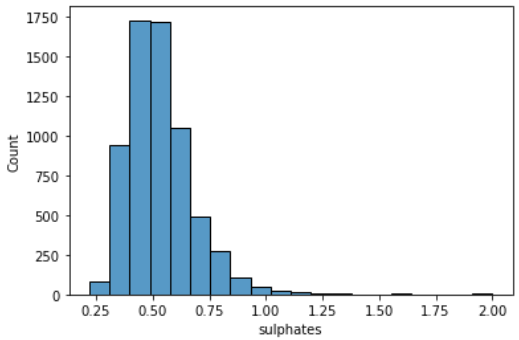
1. CRISP-DM: Data Understanding – Visual EDA
2. Univariate Visualizations
3. Bivariate Visualizations
4. Multivariate Visualizations

What is Univariate Analysis?

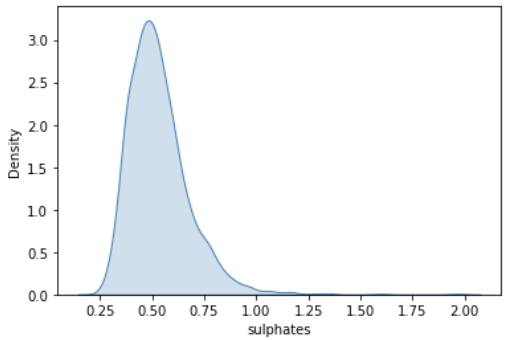
- Definition: examine one variable at a time
- Key Questions Answered
  - What is the distribution of the variable?
  - Are there any outliers?
  - Is the data skewed or normally distributed?
- Data Types of Variables
  - Numerical (continuous & discrete)
    - Use Histograms & KDE Plots
    - Use Boxplots, Violinplots
  - Categorical (nominal & ordinal)
    - Use Frequency/Count Plots
- Choosing the right visualization depends on the data type and research question
- Do not use: Pie Charts, any 3D Chart → too hard to correctly interpret for Humans

Numerical Univariate Data – Histograms & KDE Plots

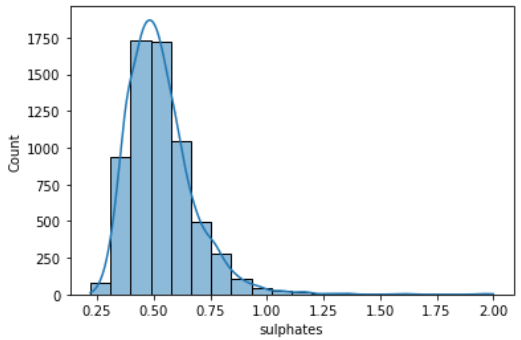
- Histogram: Groups data into bins to show frequency distribution
- Density Plots: smooth version of a histogram that estimates data density
- Use cases
  - Detect skewness, modality (uni- / bi- / multimodal distributions)
  - Understand if data follows a normal, uniform, or skewed distribution



Histogram



Density Plot

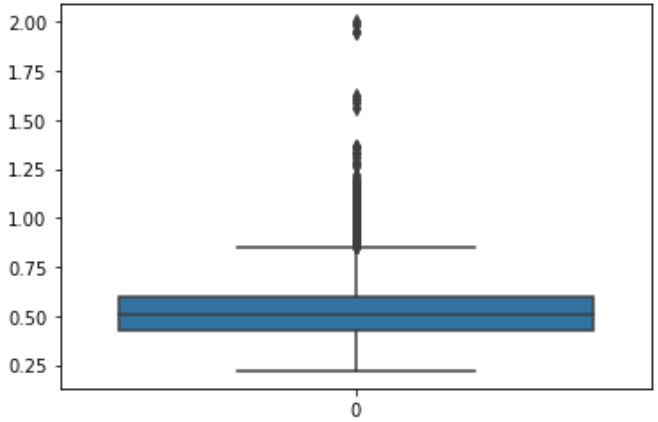


Combined Histogram and Density Plot

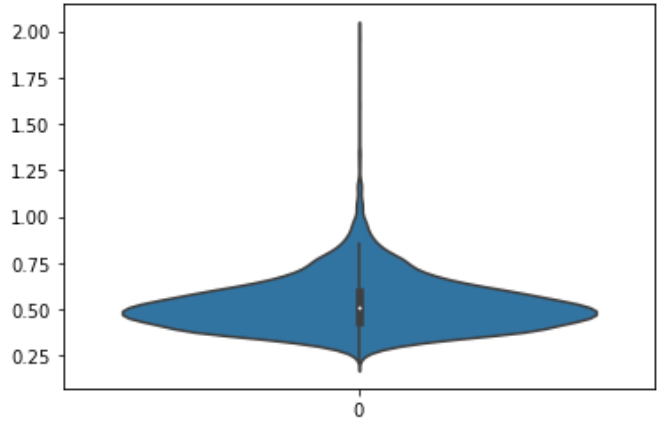
© Markus Breunig

Numerical Univariate Data – Boxplots and Violinplots

- Box Plots: Shows median, quartiles (Q1, Q3), whiskers, and outliers
- Violin Plots: Like Boxplots, but showing the KDE instead of the box
- Use cases
  - Identify outliers
  - Compare distributions across categories



Box Plot

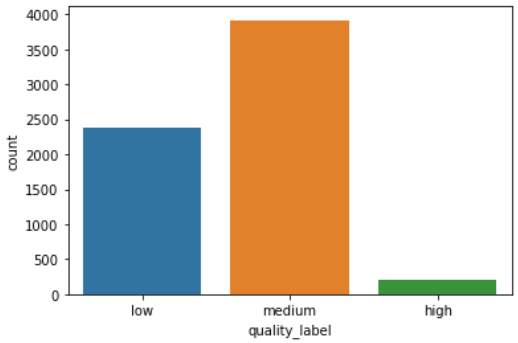


Violin Plot

© Markus Breunig

Categorical Univariate Data – Frequency/Count Plots

- Frequency/Count Plot: Shows frequency/count of categorical values
- Use cases
  - Compare category sizes in a dataset
  - Class label: evaluate if balanced or imbalanced classes



Frequency Plot

Demo

Demo

Visual EDA (Univariate)



Source: Foto von Markus Spiske auf Unsplash

Exercise

Exercise 1

Iris Data – Univariate Visualizations for EDA

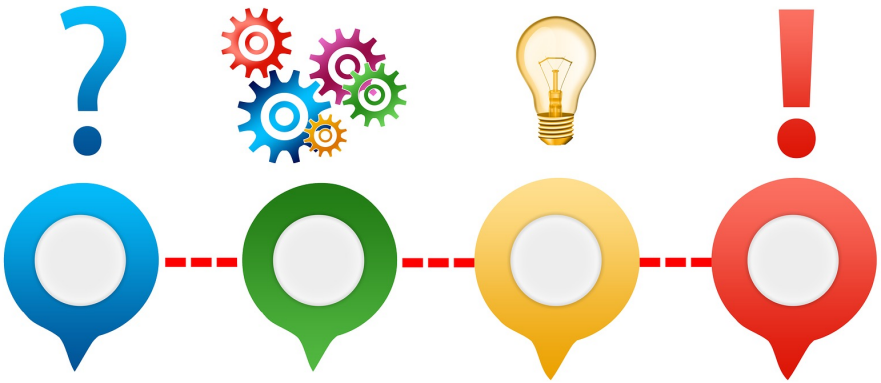


Photo by Gerd Altmann on Pixabay

Visual EDA

1. CRISP-DM: Data Understanding – Visual EDA
2. Univariate Visualizations
3. Bivariate Visualizations
4. Multivariate Visualizations

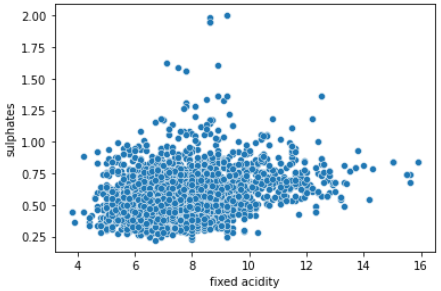
What is Bivariate Analysis?

- Definition: examine the relationship between two variables
- Key Questions Answered
  - Are two variables correlated?
  - Does one variable influence another?
  - Are there patterns or trends in the relationship?
- Types of Bivariate Relationships:
  - Numerical-Numerical
    - Use Scatter Plots / Joint Plots
    - Use Correlation Heatmaps
  - Categorical-Categorical
    - Use Multi Bar Plots
  - Categorical-Numerical
    - Use Multi Boxplots & Multi Violinplots
    - Use Strip Plots & Swarm Plots

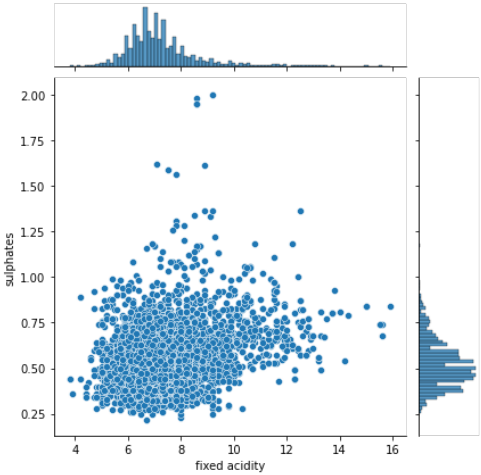
© Markus Breunig

Numerical Bivariate Data – Scatter Plots & Joint Plots

- Scatter Plot
  - Shows the relationship between two numerical variables
  - Each point represents an observation
- Joint Plot
  - Add Histograms/KDE Plots on the sides
- Use cases
  - Identify positive, negative, or no correlation
  - Detect clusters and outliers



Scatter Plot

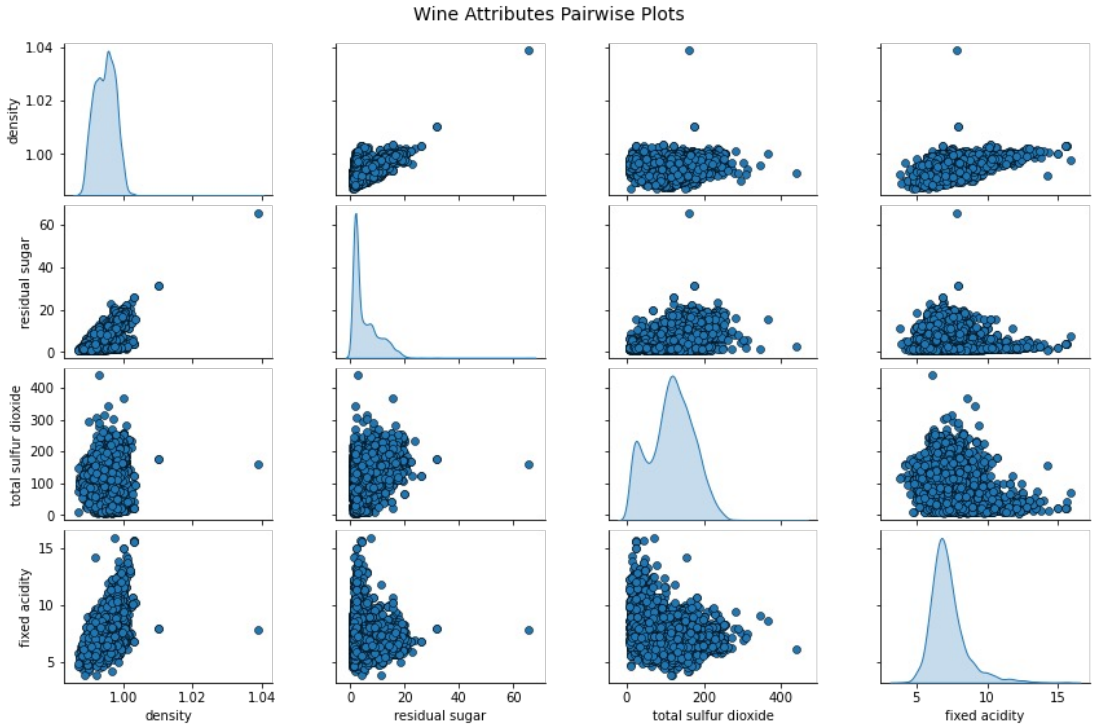


Joint Plot

© Markus Breunig

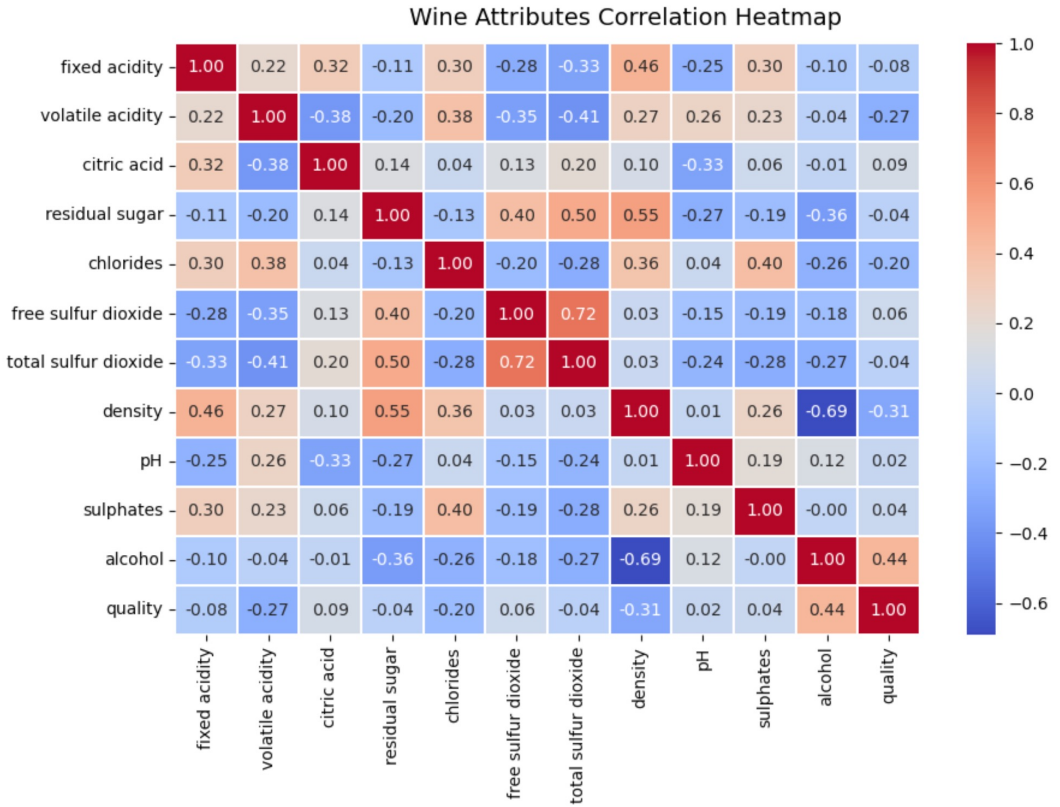
Numerical Bivariate Data – Pair-Wise Scatter Plots

- Pair-Wise Scatter Plots
  - Show all combinations & KDE Plots
- Use Cases
  - Get a quick overview for small dataset with not-too-many numerical attributes
  - Very expensive for large datasets and many numerical attributes!!



Numerical Bivariate Data – Correlation Heatmap

- Correlation Heatmap
  - Measures how strongly two numerical variables are related.
  - Most frequently based on Pearson or Spearman Correlation
  - Often coded from Blue to Red (`coolwarm`-Colour Map)
- Use cases
  - Quickly spot relationships between multiple variables
  - Helps in feature selection for machine learning



© Markus Breunig

Numerical Bivariate Data – Correlation Coefficients for Two Variables

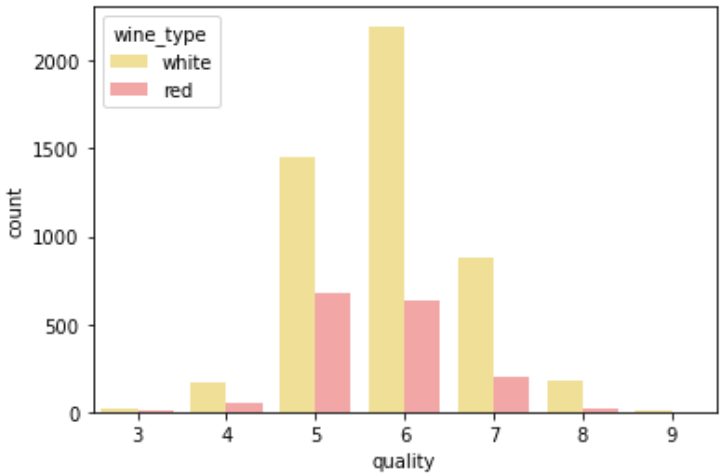
- Pearson Correlation Coefficient (r)
  - Measures how **linearly related** two variables are
  - Use it when:
    - You suspect a linear relationship
    - Data is continuous, roughly normally distributed, and free from extreme outliers
    - pandas: `df.corr(method='pearson')` (default)
- Spearman Rank Correlation Coefficient (ρ or rs)
  - Measures how well the relationship can be described by a **monotonic function**
    - Consistently increasing or decreasing, not necessarily linear
  - Use it when:
    - You expect a monotonic (but not necessarily linear) relationship
    - Data has outliers, is ordinal, or is not normally distributed
    - pandas: `df.corr(method='spearman')`

Feature	Pearson	Spearman
Type of relationship measured	Linear relationships	Monotonic relationships (increasing or decreasing)
Based on	Raw data values	Ranks of the data
Sensitivity to outliers	High	Low
Assumes	Normally distributed, continuous variables	No assumption of distribution
Range of values	[-1;1]	[-1;1]
Interpretation	1 = perfect positive linear 0 = no linear relationship -1 = perfect negative linear	Same, but for monotonic relationships

© Markus Breunig

Categorical Bivariate Data – Multi Bar Plot

- Multi Bar Plot (Multi Count Plot, Count Plot, Multi Frequency Plot)
  - Shows the frequency of categories split by another category
- Use cases
  - Compare one distribution across another distribution



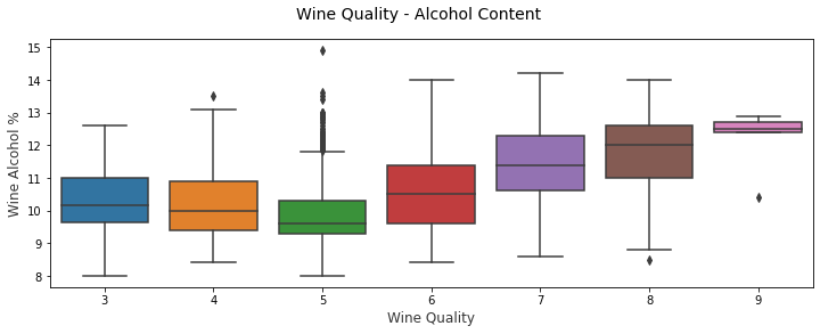
Multi Bar Plot

© Markus Breunig

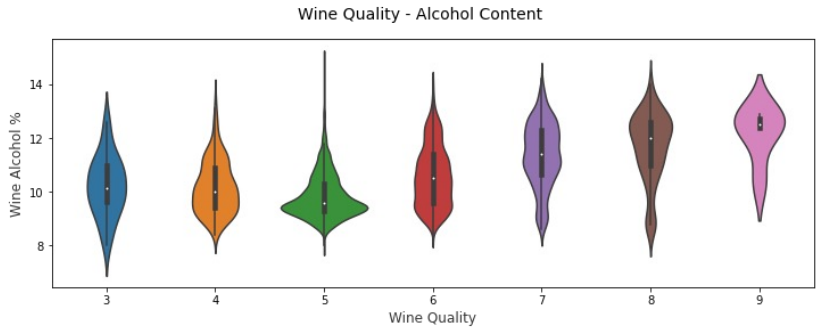
Numerical-Categorical Bivariate Data – Multi Box Plot & Multi Violin Plot

- Multi Box Plot: Shows the distribution of a numerical variable across categories
- Use cases
  - Compare distributions of numerical data across categories
  - Identify outliers

- Multi Violin Plot: Combines boxplot with KDE plot for richer insights
- Use cases
  - Identify data distribution within each category
  - Detect skewness and multimodal distributions



Multi Box Plot

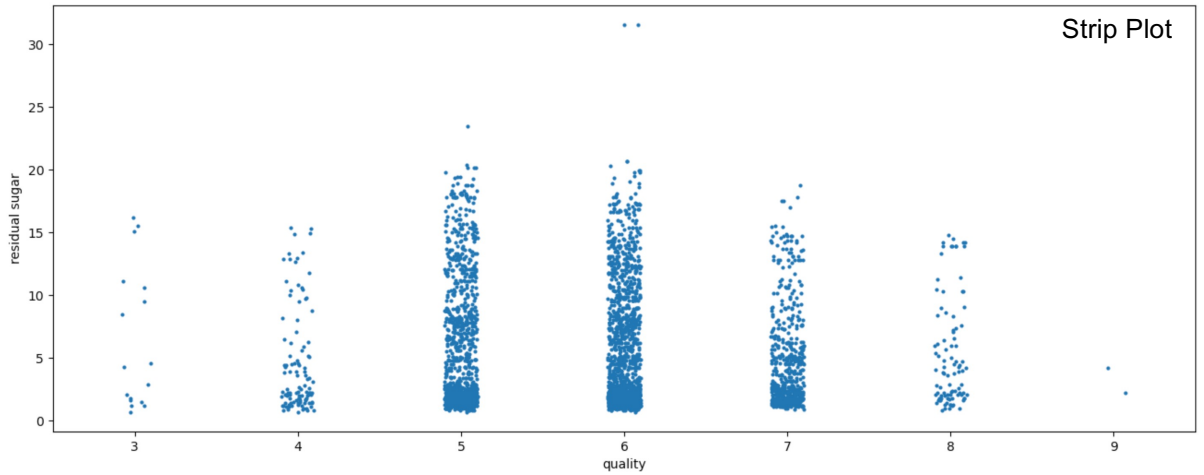


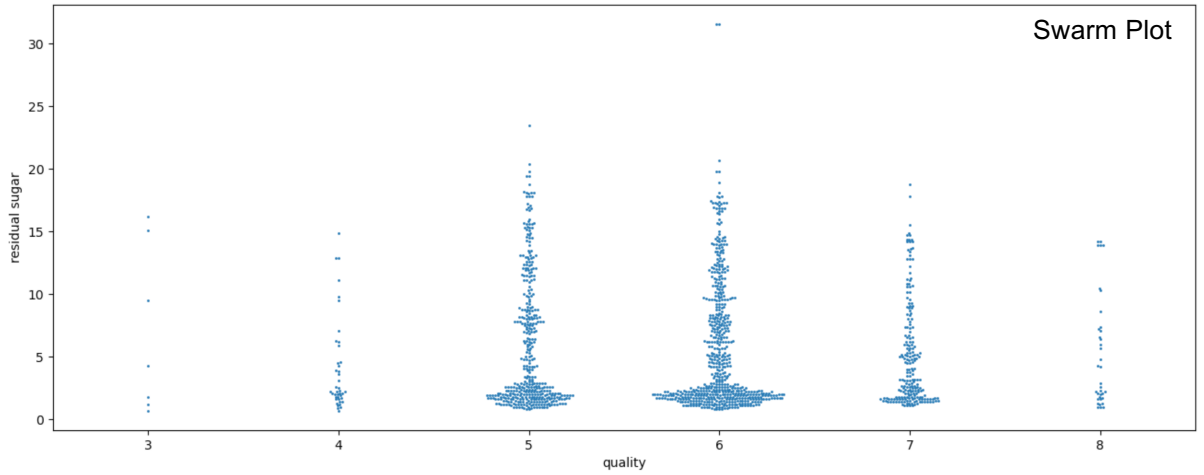
Multi Violin Plot

© Markus Breunig

Numerical-Categorical Bivariate Data – Strip Plots & Swarm Plots

- Strip Plot
  - Categorical scatterplot using jitter to reduce overplotting
- Swarm Plot
  - Same but with non-overlapping points
- Use Cases
  - Spot overlapping values and outliers
- Note: the two graphs show two different 20% random samples





© Markus Breunig

Visual EDA

1. Introduction
2. Univariate Visualizations
3. Bivariate Visualizations
4. Multivariate Visualizations

© Markus Breunig

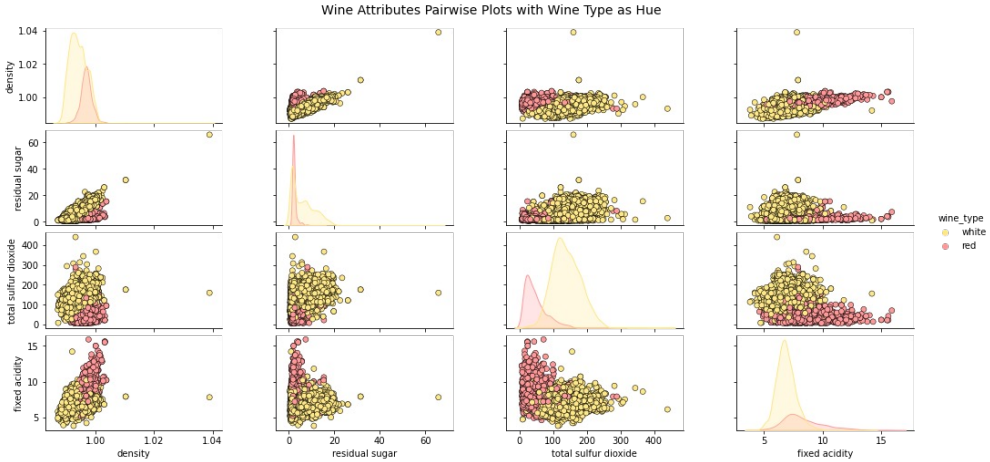
What is Multivariate Analysis?

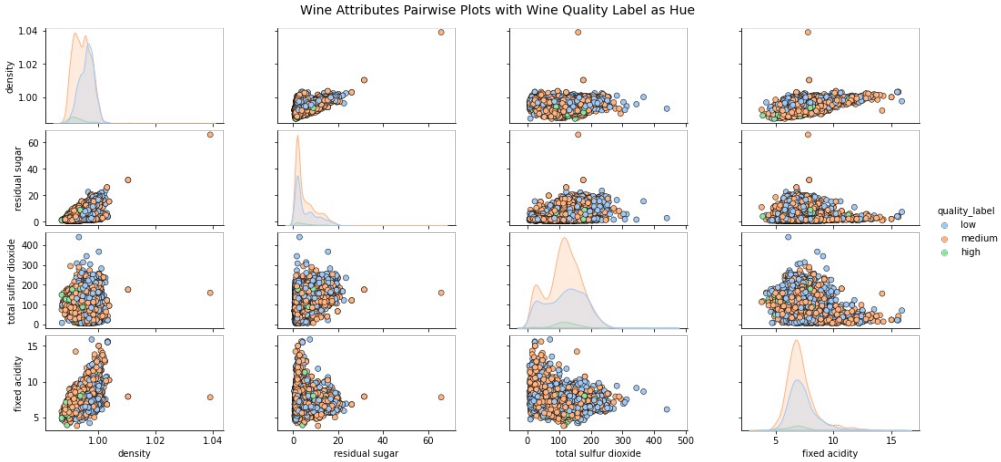
- Definition: examine the relationship between three or more variables
- Key Questions Answered
  - Are there complex patterns that bivariate analysis misses?
  - How do multiple variables interact?
  - Are there hidden relationships in the data?
- Often hard to implement and interpret!
- Approaches
  - Interactive visualizations can be helpful
    - Rotate a 3-d scatterplot
    - Time dimension as animation
  - Introduce third dimension as colour / hue / size into bivariate visualizations
    - Frequently used for the target variable in classification problems

© Markus Breunig

Multivariate – Pairwise Scatterplots with Hue as Third Dimension

- Pairwise Scatterplots with Hue as Third Dimension
  - Example 1: Hue for Wine Type
    - Ok
  - Example 2: Hue for Wine Quality
    - Hard to interpret

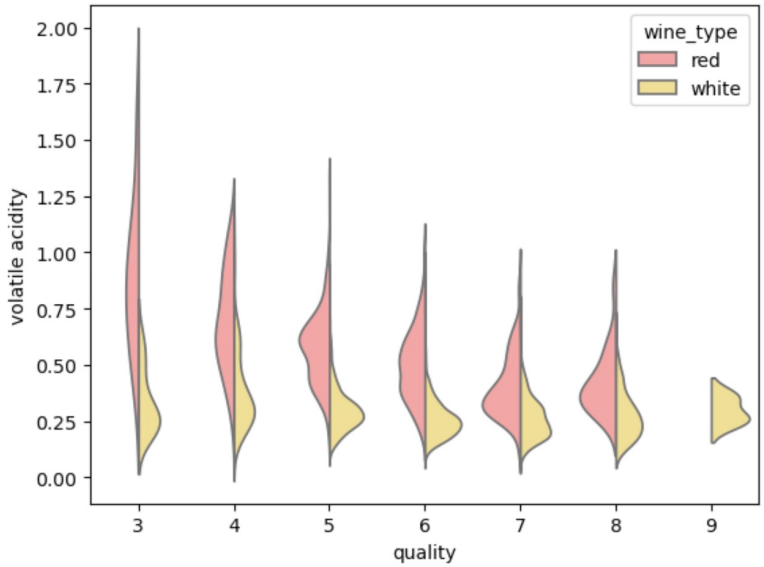




© Markus Breunig

Multivariate – Multi Violinplot with Hue for Binary Attribute

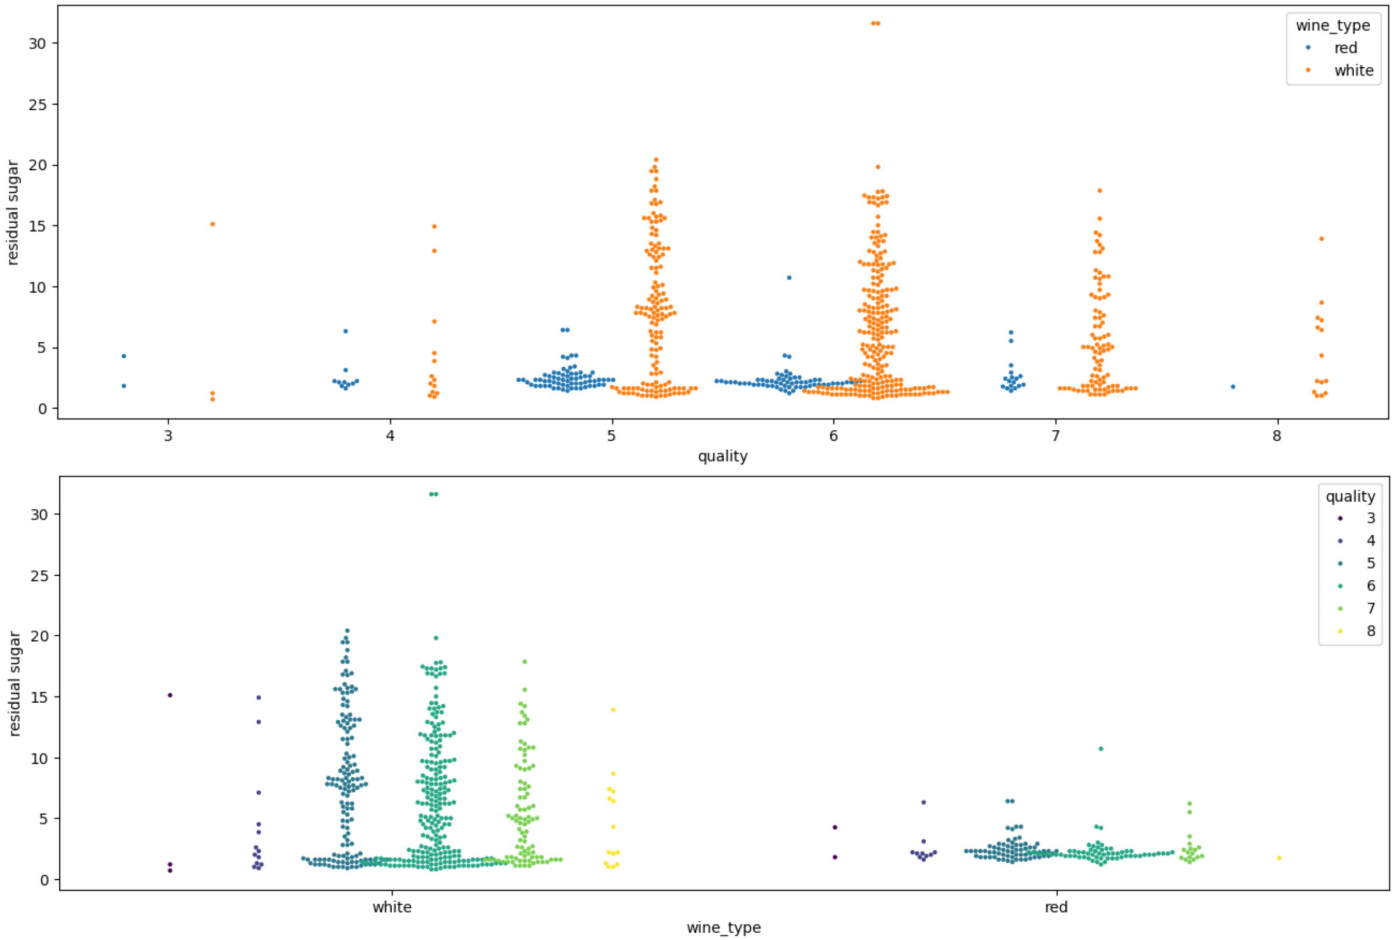
- Multi Violinplot with Hue for Binary Attribute
- Example
  - X: Quality (Categorical)
  - Y: Acidity (Numerical)
  - Hue: Wine Type (Binary)



© Markus Breunig

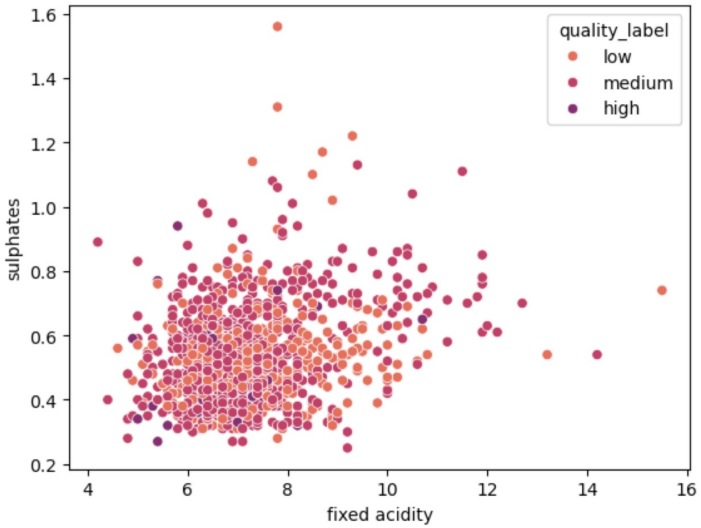
Multivariate – Swarm Plot with Hue as Third Dimension

- Swarm Plot with Hue as Third Dimension
  - Example 1
    - X: Quality (Categorical)
    - Y: Sugar (Numerical)
    - Hue: Type (Binary)
  - Example 2:
    - X: Type (Binary)
    - Y: Sugar (Numerical)
    - Hue: Quality (Cat.)
- Note: the two graphs show two different 10% random samples

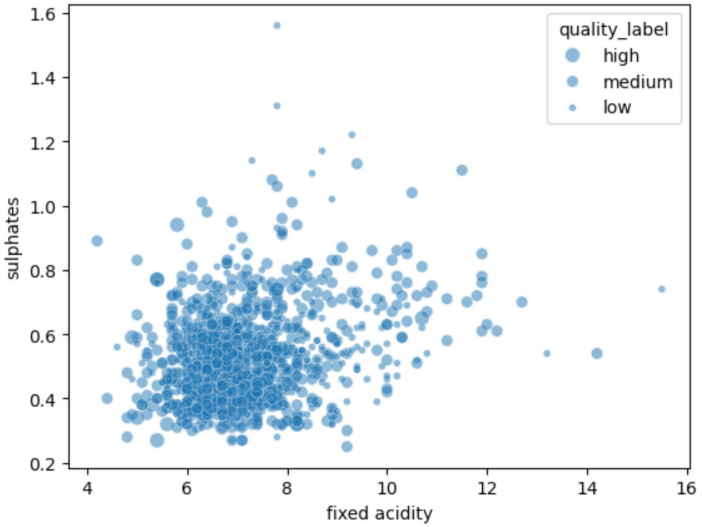


Multivariate – Scatterplot with Hue or Size as Third Dimension

- Scatterplot with Hue as Third Dimension
  - X: Numerical (fixed acidity)
  - Y: Numerical (sulpathes)
  - Hue: Categorical (quality_label)



- Scatterplot with Size as Third Dimension (Bubble Chart)
  - X: Numerical (fixed acidity)
  - Y: Numerical (sulpathes)
  - Size: Numerical or Categorical (here: categorical, quality_label)



Note: the two graphs show the same 20% random sample

© Markus Breunig

Demo

Demo

Visual EDA (Bi- and Multivariate)



Source: Foto von Markus Spiske auf Unsplash

Common Mistakes and Best Practices

- Common Mistakes in Visual EDA
  - Ignoring categorical variables when plotting numerical data
  - Using the wrong type of visualization (e.g., pie charts)
  - Not labelling axes or providing enough context for plots
- Some Best Practices for Effective Data Visualization
  - Choose the right visualization for your data type
  - Use appropriate labels, legends, and titles
  - Avoid misleading scales or truncation
  - Colour-code data effectively (e.g., divergent colours for negative vs. positive values)
  - Check for overplotting and adjust transparency (alpha parameter)

© Markus Breunig

Exercise

Exercise 2

Iris Data – Bivariate and Multivariate Visualizations for EDA

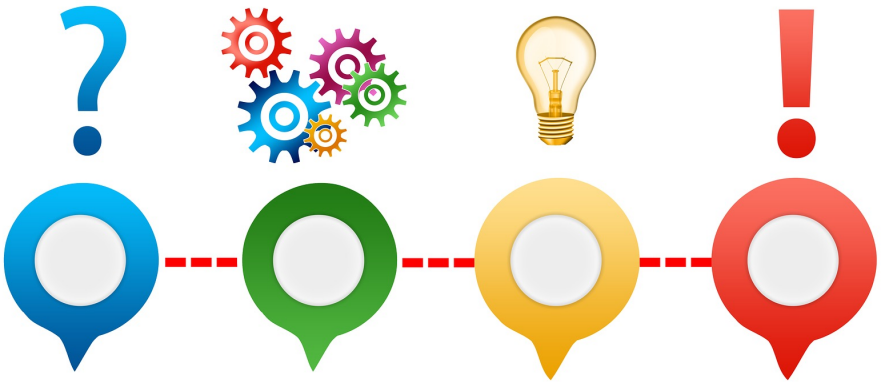


Photo by Gerd Altmann on Pixabay

Key Takeaways

- Visual Analysis of the data is key part of EDA
- Start with univariate visualizations of all features (and the target, if applicable)
  - Start with a Histogram/Density Plot/Frequency Plot
  - Look for Outliers in the Box Plots
- Continue with bivariate/multivariate visualizations
  - Pair-wise scatter plots
  - Colour code/hue code (the target), if applicable
- Use easy-to-interpret and easy-to-communicate visualizations
- Focus – explore deep but do not get carried away



Image by Gordon Johnson on pixabay