



Chapter 2 – Relational models and relational algebra


Databases lectures

Dr Kai Höfig

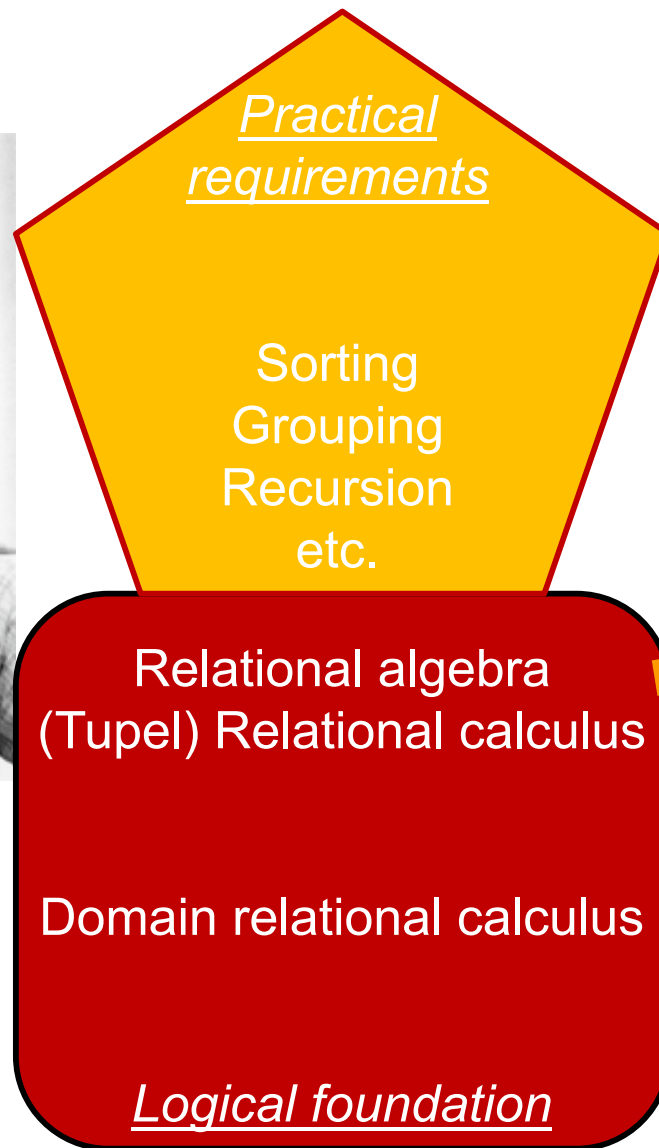


Relational algebra, calculus and languages

Theory




In the 1960s and 1970s, Edgar Codd created the relational model that forms the basis for relational databases, which are still a standard of database technology today.



Without understanding the logical foundation, we interact with an unknown communication partner using the database languages

Practice





Relational algebra and SQL

- ♦ **Relational algebra** is the foundation for the **tuple relational calculus**, which is implemented by the **Sequence Query Language** (SQL)
 - Relational calculus is **declarative**, i.e. the calculation sequence is not visible
 - Easier to use than procedural instructions – user does not need to know how the DBMS calculates the result.
 - DBMS has freedom with processing, so it can choose one that is as efficient as possible.
 - Relational algebra is **procedural**, i.e. it specifies the processing sequence (“from inside to outside”).
 - There are laws for the transformation of algebra expressions.
 - DBMS translates the SQL query (=tuple relational calculus) into an algebra expression, optimises it (query optimiser) and executes it.
 - This “execution plan” is used by the DB administrator to optimise the DB (e.g. by means of index structures).



The relational data model

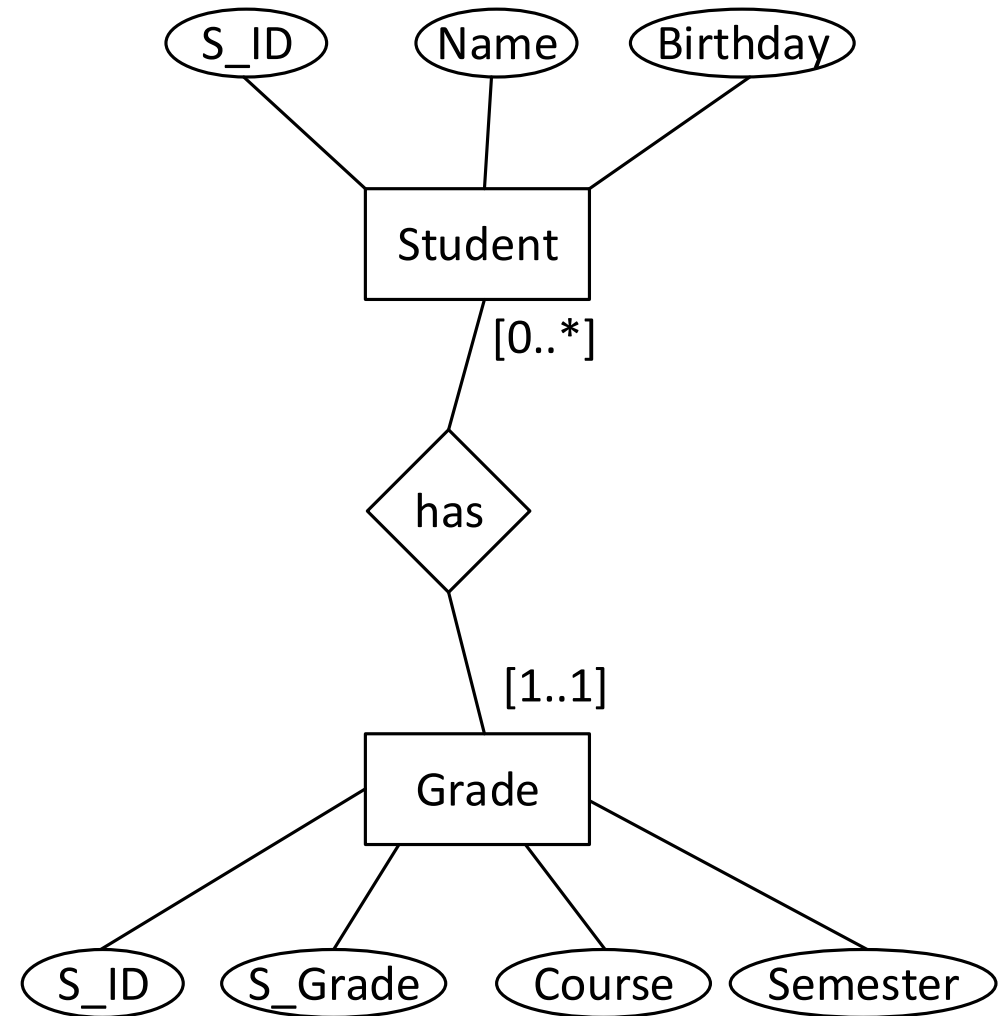
- ◆ **Structure** of the data:
Data is stored in relations (tables)
- ◆ **Operations** on the data: 2 alternatives
 1. Relational calculus (implemented in SQL)
 2. Relational algebra (base for SQL)
- ◆ **Integrity constraints**: The most important:
 1. Key constraints
 2. Referential integrity = foreign key constraints
 3. Domain constraints (restrictions on values allowed for an attribute)
 - Integrity constraints can be formulated as conditions in relational algebra, relational calculus or SQL



Running example for this lecture

Students:	<u>S_ID</u>	Name	Birthday
	1	Max	14.05.2001
	2	Miri	17.04.2000
	3	Sarah	18.09.2002
	4	Ben	07.06.2002

Grade:	<u>S_ID</u>	<u>S_Grade</u>	Course	Semester
	1	1.0	DB	WS23
	2	2.3	DB	WS23
	2	1.3	OOP	SS23
	3	4.0	OOP	SS22
	4	3.3	DB	WS23






Representation of relations and terms

◆ Representation

- First row: **relation schema**
- Other entries in the table: **relation**
- One row of the table: **tuple**
- A column heading: **attribute**
- One entry: **attribute value**





Definition of the relational data model (1)

- ◆ In the relational data model, we only have the relational schema for structural modelling.

- ◆ A **relational schema** with the identifier R and Attributes A

$$R = (A_1, \dots, A_n)$$

- ◆ A **domain function** assigning value ranges to the attributes

$$dom : \{A_1, \dots, A_k\} \rightarrow \{D_1, \dots, D_s\}, s \geq 1$$

- ◆ A **relational database schema** containing the relational schema

$$S = \{R_1(A_1, \dots, A_i), \dots, R_m(A_j, \dots, A_k)\}$$

$$\begin{aligned} \text{GradingSystem} &= \{Students(S_ID, Name, Birthday), \\ &= Grade(S_ID, S_Grade, Course, Semester)\} \end{aligned}$$

$$dom : \{S_ID, Name, Birthday, S_Grade, Course, Semester\} \rightarrow int, varchar$$



Definition of the relational data model (2)

- ♦ A **relation** r of a relational schema R is a finite set of tuples and is also referred to as **base relation**.

$$r(R) = t_0, \dots, t_n$$

- ♦ A **database** via a database schema S is a set of relations.

$$d = r_0, \dots, r_n$$

$$\begin{aligned} r_{Students} &= t_0, t_1, t_2, t_3 \\ d_{GradingSystem} &= r_{Students}, r_{Grades} \\ t_0(Name) &= Max \end{aligned}$$

Students:	<u>S_ID</u>	<u>Name</u>	<u>Birthday</u>
	1	Max	14.05.2001
	2	Miri	17.04.2000
	3	Sarah	18.09.2002
	4	Ben	07.06.2002



Important integrity constraints in RM: **Key constraint**

- ◆ One or more attributes fulfil (at a given point in time) the **key property**, if there are no two tuples in that relation with the same value of those attributes. We also say, that one or more attributes **identify** uniquely stored tuples.

Key constraint: There are no 2 tuples in *Students* with the same value of *S_ID*. From the domain, we know that this is a candidate for a key.

Students:			
<u>S_ID</u>	Name	Birthday	
1	Max	14.05.2001	
2	Miri	17.04.2000	
3	Sarah	18.09.2002	
4	Ben	07.06.2002	

- ◆ Note

- A key is domain knowledge
- Attribute combinations can also be keys
- A table can have multiple keys
- When the DB is designed, one key is selected as the **primary key**
- Keys are generally marked by underlining

Key constraint: There are no 2 tuples in *Students* with the same value of *Name*. Despite the key constraint is currently fulfilled, we know that is cannot be a key.



Exercise

- ◆ Which single attributes fulfil the **key attribute**?
- ◆ Which attribute **combinations** fulfil the key attribute?
- ◆ What attributes are **key candidates**?
- ◆ What are **primary keys**?

Students:	<u>S_ID</u>	<u>Name</u>	<u>Birthday</u>
	1	Max	14.05.2001
	2	Miri	17.04.2000
	3	Sarah	18.09.2002
	4	Ben	07.06.2002

Grade:	<u>S_ID</u>	<u>S_Grade</u>	<u>Course</u>	<u>Semester</u>
	1	1.0	DB	WS23
	2	2.3	DB	WS23
	2	1.3	OOP	SS23
	3	4.0	OOP	SS22
	4	3.3	DB	WS23



Important integrity constraints in RM: **Referential Integrity**

- ◆ Two sets of attributes are of **referential integrity**, if one set of attribute values is a subset of the other set of attribute values.

Students:	<u>S_ID</u>	Name	Birthday
	1	Max	14.05.2001
	2	Miri	17.04.2000
	3	Sarah	18.09.2002
	4	Ben	07.06.2002

Grade:	<u>S_ID</u>	<u>S_Grade</u>	Course	Semester
	1	1.0	DB	WS23
	2	2.3	DB	WS23
	2	1.3	OOP	SS23
	3	4.0	OOP	SS22
	4	3.3	DB	WS23

- ◆ Here, the values of S_ID in `Grade` is a subset of S_ID in `Students`.
- ◆ Typically, a foreign key in one relation is a subset of a key in another relation



Structure is set!

Students:

<u>S_ID</u>	Name	Birthday
1	Max	14.05.2001
2	Miri	17.04.2000
3	Sarah	18.09.2002
4	Ben	07.06.2002

Grade:

<u>S_ID</u>	<u>S_Grade</u>	Course	Semester
1	1.0	DB	WS23
2	2.3	DB	WS23
2	1.3	OOP	SS23
3	4.0	OOP	SS22
4	3.3	DB	WS23

Attribute

Relational
schema

Tuple

Key

Referential
integrity

Attribute
Value

Relation



But how do I generate relations from base-relations?

Students: S_ID Name Birthday

1	Max	14.05.2001
2	Miri	17.04.2000
3	Sarah	18.09.2002
4	Ben	07.06.2002

Grade: S_ID S_Grade Course Semester

1	1.0	DB	WS23
2	2.3	DB	WS23
2	1.3	OOP	SS23
3	4.0	OOP	SS22
4	3.3	DB	WS23

„I can recommend all students as tutors, that have a grade better than 2.0“



Query operations on tables

- ◆ **Relational algebra**: set of **basic operations** on relations to compute new (resulting) relations from base relations
 - can be combined in any way
 - thereby create an algebra for “calculating with tables”
 - Provides the background for the calculation of declarative languages like SQL
- ◆ Revision from mathematics:
algebra = value range + operations defined on these
- ◆ Here
 - value range = contents of the database = tables
 - operations = functions for calculating new tables



Relational algebra: overview

- ◆ Three main operations: Selection, Projection, Cross Product

A	B	C
a1	b1	c1
a2	b1	c2
a3	b2	c1
a4	b3	c3
a5	b1	c1

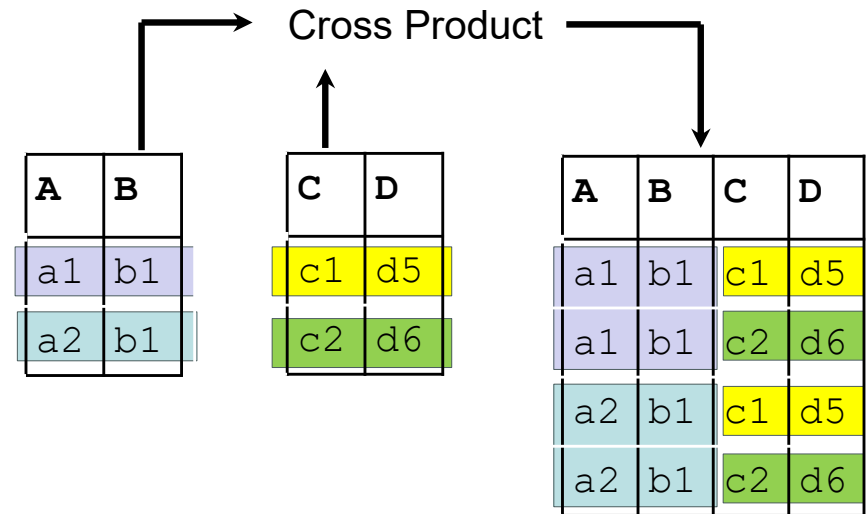
Projection

π

A	B	C
a1	b1	c1
a2	b1	c2
a3	b2	c1
a4	b3	c3
a5	b1	c1

Selection

σ



Cross Product

\times



Projection π Definition

- ◆ **Projection π (\underline{P} i)**: selection of columns by specifying an attribute list

- ◆ **Syntax** $\pi_{\langle AttributeList \rangle}(Relation)$


- ◆ **Example**
- | R | <table border="1"><thead><tr><th>A</th><th>B</th></tr></thead><tbody><tr><td>1</td><td>2</td></tr><tr><td>1</td><td>3</td></tr><tr><td>3</td><td>4</td></tr></tbody></table> | A | B | 1 | 2 | 1 | 3 | 3 | 4 | $\pi_A(R)$ | <table border="1"><thead><tr><th>A</th></tr></thead><tbody><tr><td>1</td></tr><tr><td>3</td></tr></tbody></table> | A | 1 | 3 |
|-----|--|---|---|---|---|---|---|---|---|------------|---|---|---|---|
| A | B | | | | | | | | | | | | | |
| 1 | 2 | | | | | | | | | | | | | |
| 1 | 3 | | | | | | | | | | | | | |
| 3 | 4 | | | | | | | | | | | | | |
| A | | | | | | | | | | | | | | |
| 1 | | | | | | | | | | | | | | |
| 3 | | | | | | | | | | | | | | |

- ◆ The projection removes duplicate tuples
- ◆ **Set semantics**, \rightarrow only in relational algebra!



Projection: Augmentation

- ◆ For $A \subseteq B \subseteq R : \pi_A(\pi_B(R)) = \pi_A(R)$



- ◆ Optimization: when projecting on attributes and then projecting on a subset, we can directly project on that subset



Selection σ Definition

- ◆ Selection σ (Sigma): selection of rows of a table based on a selection predicate

- ◆ Syntax $\sigma_{\langle Constraint \rangle}(Relation)$

- ◆ Example


R	<table><tr><th>A</th><th>B</th></tr><tr><td>1</td><td>2</td></tr><tr><td>1</td><td>3</td></tr><tr><td>3</td><td>4</td></tr></table>	A	B	1	2	1	3	3	4	$\sigma_{A=1}(R)$	<table><tr><th>A</th><th>B</th></tr><tr><td>1</td><td>2</td></tr><tr><td>1</td><td>3</td></tr></table>	A	B	1	2	1	3
A	B																
1	2																
1	3																
3	4																
A	B																
1	2																
1	3																

- ◆ Conditions: Attributes and constants using $=$, \neq , \leq , $<$, \geq or $>$, linking and combining of multiple conditions using \wedge , \vee or \neg and $(,)$.



Selection: Commutativity

$$\sigma_{A=a}(\sigma_{B=b}(R)) = \sigma_{B=b}(\sigma_{A=a}(R))$$

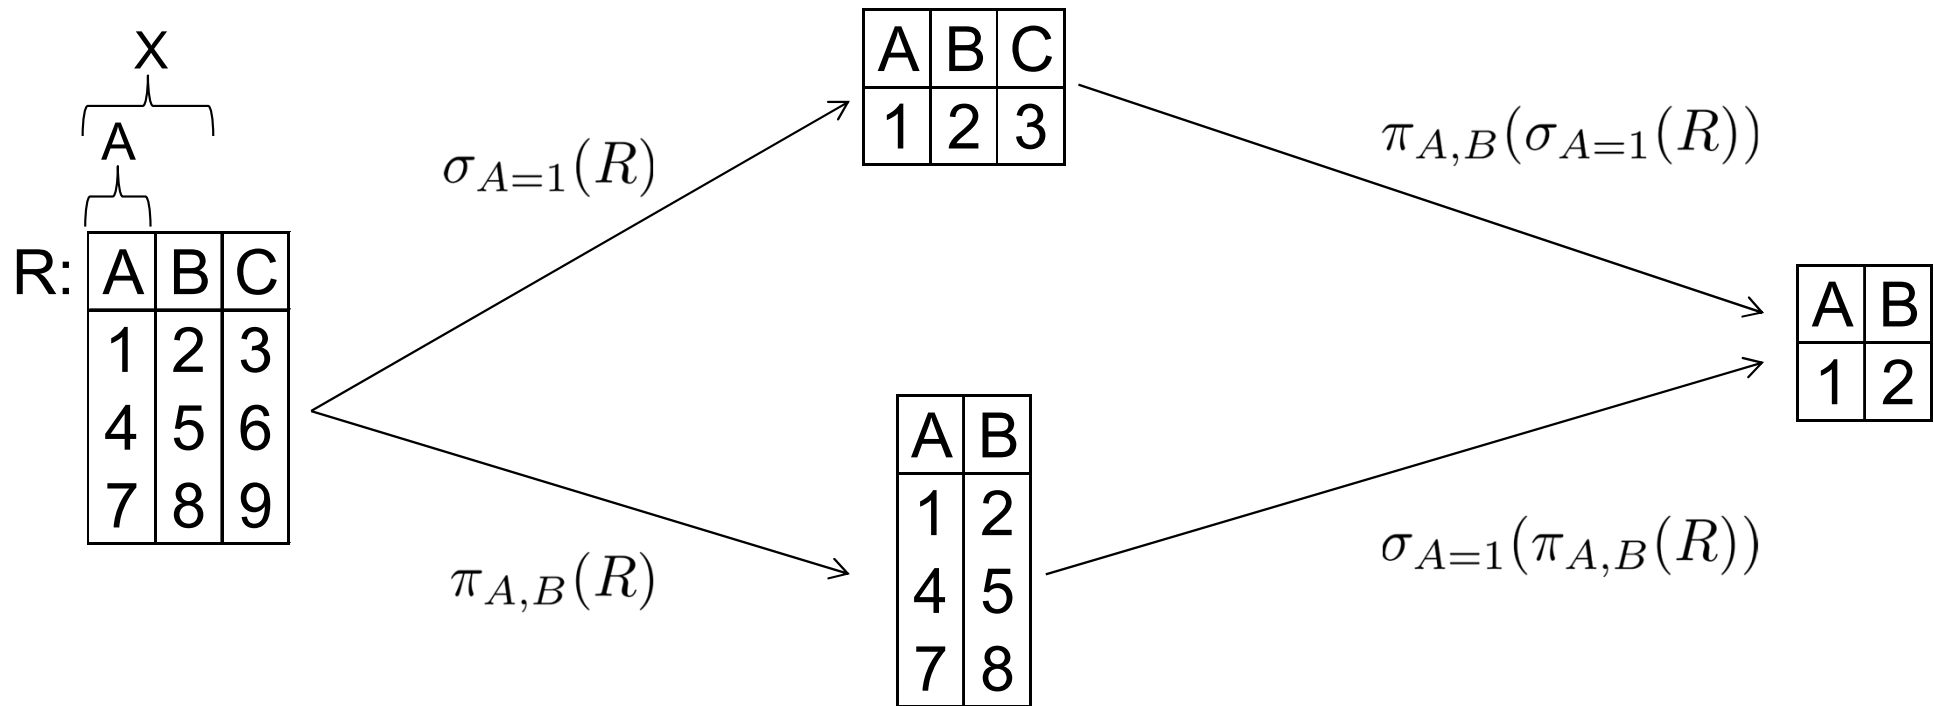


- ♦ Optimization: selection is commutative and sometimes changing the order of execution can be more efficient



Selection: Transposition

- ◆ If $A \in X, X \subseteq R : \pi_X(\sigma_{A=a}(R)) = \sigma_{A=a}(\pi_X(R))$



- ◆ Optimization: Sometimes first select and then project can be more efficient



Cross product \times

- ◆ **Cross product \times (Cartesian product, cross join):** links two tables by combining each tuple of the first with each tuple of the second.
 - Be careful: result for tables with n or m tuples has $n*m$ tuples!
- ◆ **Syntax** $\langle \text{Relation1} \rangle \times \langle \text{Relation2} \rangle$
- ◆ **Semantics** $R \times S := \{x_1, \dots, x_n, \dots, x_{n+m} \mid x_1, \dots, x_n \in r, x_n, \dots, x_{n+m} \in s\}$
- ◆ **Example:**

R:

A	B
1	2
3	4

S:

C	D
5	6
7	8

$R \times S$:

A	B	C	D
1	2	5	6
1	2	7	8
3	4	5	6
3	4	7	8

Attributes with the same name are uniquely identified by prefixing the relations name



But how do I generate relations from base-relations?

Students: S_ID	Name	Birthday
1	Max	14.05.2001
2	Miri	17.04.2000
3	Sarah	18.09.2002
4	Ben	07.06.2002

Conceptual
Schema

Grade: S_ID	S_Grade	Course	Semester
1	1.0	DB	WS23
2	2.3	DB	WS23
2	1.3	OOP	SS23
3	4.0	OOP	SS22
4	3.3	DB	WS23

„I can recommend all students as tutors, that have a grade better than 2.0“

σ, π, \times

External
Schema

Let us start with

Students \times *Grade*



Another example using renaming

„I need a list of persons, their children and grandchildren“

R:	Person	Child
	Charles the Great Louis the Pious Louis the Pious Lothar the I	Louis the Pious Lothar the I Charles the Bald Louis the II

σ, π, \times



Set operations: union

- ◆ **Union** of two relations: collects the tuple sets of two relations under a common schema
- ◆ Attribute sets of both relations must be **identical**, we care for attribute names, we do not care for attribute order.
- ◆ Example:

R:

A	B	C
1	2	3
4	5	6

S:

A	B	C
4	5	6
7	8	9

$R \cup S$:

A	B	C
1	2	3
4	5	6
7	8	9



Example Union

Students:			Grade:			
S_ID	Name	Birthday	S_ID	S_Grade	Course	Semester
1	Max	14.05.2001	1	1.0	DB	WS23
2	Miri	17.04.2000	2	2.3	DB	WS23
3	Sarah	18.09.2002	2	1.3	OOP	SS23
4	Ben	07.06.2002	3	4.0	OOP	SS22
			4	3.3	DB	WS23

„All student IDs that did pass DB better than 2.0 or are born before 2001“



Set operations: difference

- ◆ **Difference** eliminates the tuples from the first relation that also occur in the second relation
- ◆ Attribute sets of both relations must be identical
- ◆ Example:

R:

A	B	C
1	2	3
4	5	6

S:

A	B	C
4	5	6
7	8	9

R - S:

A	B	C
1	2	3



Set operations: intersection

- ♦ **Intersection** returns the tuples that are common to both relations
- ♦ Attribute sets of both relations must be identical
- ♦

- ♦ Example:

R:	A	B	C	S:	A	B	C	$R \cap S$:	A	B	C
	1	2	3		4	5	6		4	5	6
	4	5	6		7	8	9				

- ♦ Intersection can be replaced by difference

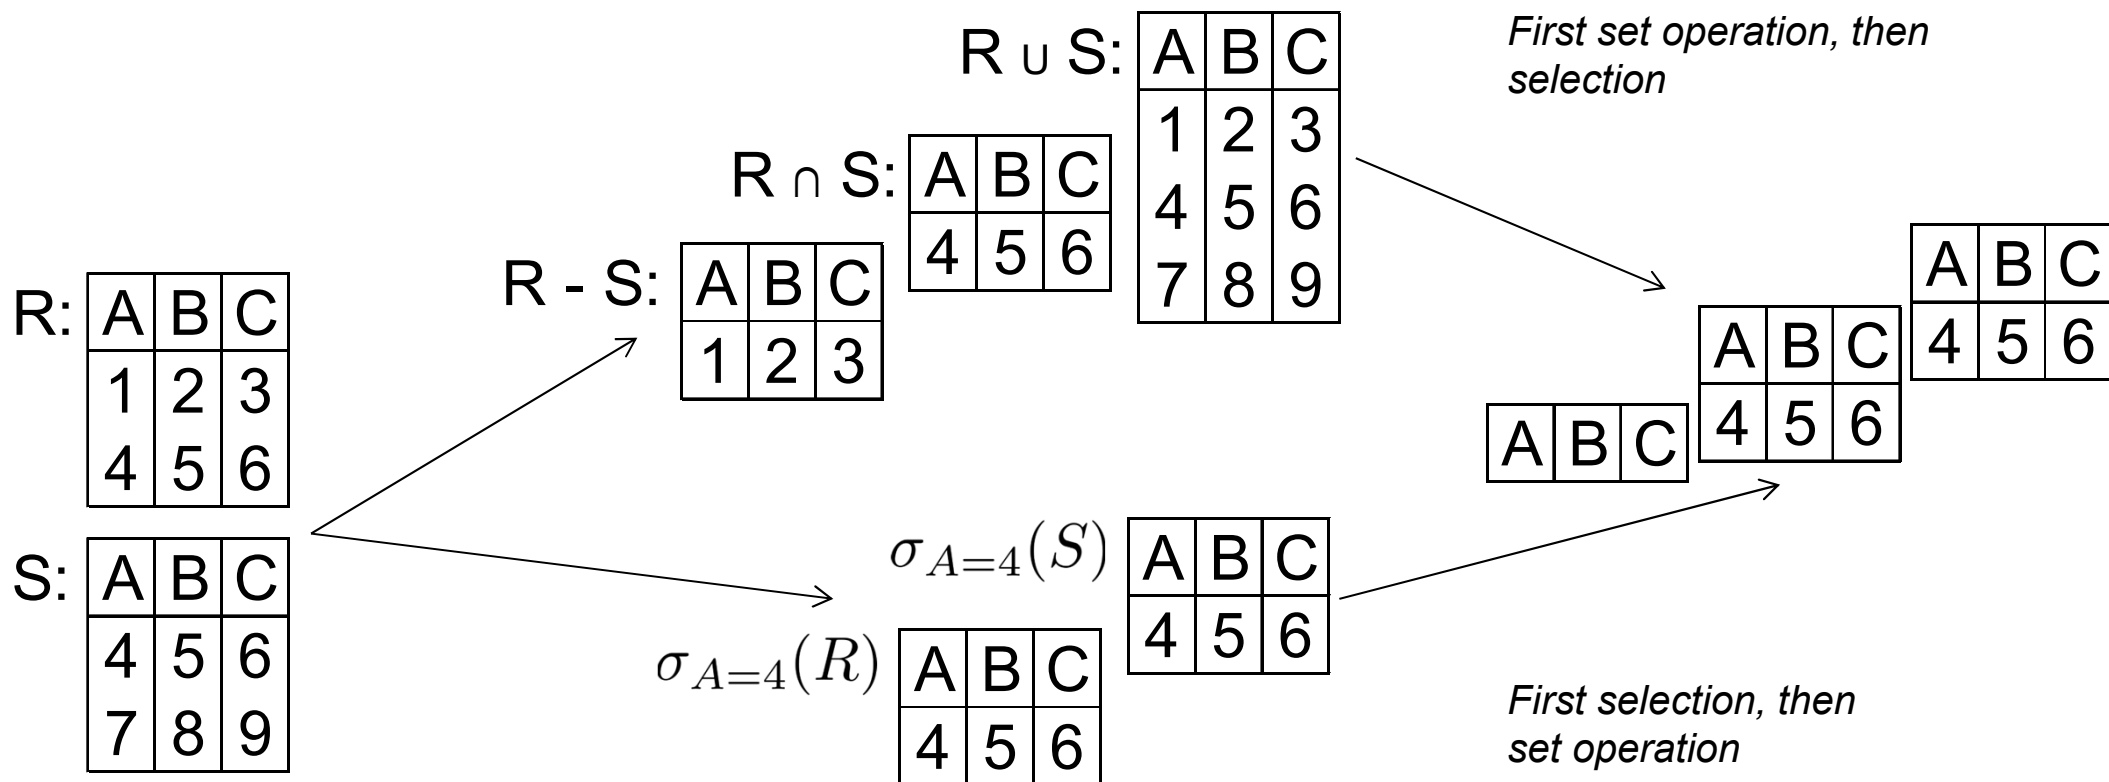
$$R \cap S = R - (R - S)$$



Set operations: Laws for transformation

- ◆ **Distributivity** regarding \cap , \cup , $-$

$$\sigma_{A=4}(R \cup S) = \sigma_{A=4}(R) \cup \sigma_{A=4}(S)$$





Relational algebra

- ◆ **Hide columns:** Projection π
- ◆ **Search for rows:** Selection σ
- ◆ **Link tables:** cross product \times
- ◆ **Unify tables:** Union \cup
- ◆ **Subtract tables from each other:** Difference and intersection \cap



Application example

„All bicycles that are produced in DE and 22" in size.“

Bike (B)

BName	BManufacturer	BSize
Stereo 150	Cube	22
Balance bike	Puky	8

Manufacturer (M)

MName	MCountry
Cube	49
Puky	49
Yeti Cycles	1

Country (C)

CCode	CName
49	DE
1	USA



Independence and completeness

- ◆ A query language is called **relationally complete** if every relational algebra operation in the language can be executed by (one or more) commands
- ◆ There is a minimal set of operations within the relational algebra from which all other operations can be composed: $\Omega = \pi, \sigma, \times, \cup$ and -
 - Ω is **independent**: no operator can be omitted without losing completeness
 - Other independent, complete sets: replace \times with \bowtie and β
- ◆ Thus: it is sufficient to show that all operations from Ω can be expressed in a query language to show that this is relationally complete
- ◆ SQL is relationally complete!



Discussion

- ◆ Which structural elements of the relational data model are we familiar with?
- ◆ What relational algebra operations are there?
- ◆ Which set of relational algebra operations is relationally complete and independent? Why is this important?
- ◆ Which integrity constraints do we know from the relational data model?